

PredictFlightOnTimePerformance

November 6, 2019

```
[1]: import pandas as pd
      # Read flight data from spreadsheet to 'pandas' model
      # read_csv returns a DataFrame (two-D data structure with labeled axes)
      flight_data = pd.read_csv('ONTIME_FLIGHT_DATA.csv')
      flight_data.head()
```

```
[1]:  YEAR  MONTH  DAY_OF_MONTH  DAY_OF_WEEK  ORIGIN_AIRPORT_ID  \
0   2019      1           19           6           13487
1   2019      1           20           7           13487
2   2019      1           21           1           13487
3   2019      1           22           2           13487
4   2019      1           23           3           13487

      ORIGIN_AIRPORT_SEQ_ID  ORIGIN_CITY_MARKET_ID  ORIGIN  DEST_AIRPORT_ID  \
0                1348702                31650    MSP           11193
1                1348702                31650    MSP           11193
2                1348702                31650    MSP           11193
3                1348702                31650    MSP           11193
4                1348702                31650    MSP           11193

      DEST_AIRPORT_SEQ_ID  DEST_CITY_MARKET_ID  DEST  CRS_DEP_TIME  ARR_DELAY_NEW  \
0                1119302                33105   CVG           1556           0.0
1                1119302                33105   CVG           1556           0.0
2                1119302                33105   CVG           1556           0.0
3                1119302                33105   CVG           1556           0.0
4                1119302                33105   CVG           1556           0.0

      ARR_DEL15  Unnamed: 15
0           0.0          NaN
1           0.0          NaN
2           0.0          NaN
3           0.0          NaN
4           0.0          NaN
```

```
[2]: print(type(flight_data))
```

```
<class 'pandas.core.frame.DataFrame'>
```

1 Clean Data

```
[3]: # before construct a model, need to clean data to remove null values
flight_data.isnull().values.any()
```

```
[3]: True
```

```
[4]: flight_data.isnull().sum()
```

```
[4]: YEAR                0
MONTH                0
DAY_OF_MONTH        0
DAY_OF_WEEK         0
ORIGIN_AIRPORT_ID   0
ORIGIN_AIRPORT_SEQ_ID 0
ORIGIN_CITY_MARKET_ID 0
ORIGIN              0
DEST_AIRPORT_ID     0
DEST_AIRPORT_SEQ_ID 0
DEST_CITY_MARKET_ID 0
DEST                0
CRS_DEP_TIME        0
ARR_DELAY_NEW       21000
ARR_DEL15           21000
Unnamed: 15         638649
dtype: int64
```

```
[5]: flight_data = flight_data.drop('Unnamed: 15', axis=1)
```

```
[6]: flight_data.isnull().sum()
```

```
[6]: YEAR                0
MONTH                0
DAY_OF_MONTH        0
DAY_OF_WEEK         0
ORIGIN_AIRPORT_ID   0
ORIGIN_AIRPORT_SEQ_ID 0
ORIGIN_CITY_MARKET_ID 0
ORIGIN              0
DEST_AIRPORT_ID     0
DEST_AIRPORT_SEQ_ID 0
DEST_CITY_MARKET_ID 0
DEST                0
CRS_DEP_TIME        0
ARR_DELAY_NEW       21000
ARR_DEL15           21000
dtype: int64
```

```
[7]: # CRS_ARR_TIME      Scheduled arrival time; ARR_DEL15      0=Arrived less
      →than 15 minutes late, 1=Arrived 15 minutes or more late
```

```
flight_data = flight_data[['MONTH', 'DAY_OF_MONTH', 'DAY_OF_WEEK', 'ORIGIN', 'DEST', 'CRS_DEP_TIME', 'ARR_DEL15']]
```

```
[8]: flight_data.isnull().sum()
```

```
[8]: MONTH          0
     DAY_OF_MONTH    0
     DAY_OF_WEEK     0
     ORIGIN          0
     DEST           0
     CRS_DEP_TIME    0
     ARR_DEL15      21000
     dtype: int64
```

```
[9]: flight_data.shape
```

```
[9]: (638649, 7)
```

```
[10]: flight_data = flight_data.fillna({'ARR_DEL15': 1}) # fill NA values with '1' in ARR_DEL15 column
```

```
[11]: import numpy as np
     # In order to avoid overfitting, divide CRS_DEP_TIME:scheduled arrival time by 100 because it matters more if flight is delayed by hours rather than by minutes
     flight_data['CRS_DEP_TIME'] = flight_data['CRS_DEP_TIME'].div(100).apply(np.floor)
```

```
[12]: flight_data.head()
```

```
[12]:   MONTH  DAY_OF_MONTH  DAY_OF_WEEK  ORIGIN  DEST  CRS_DEP_TIME  ARR_DEL15
0      1           19           6    MSP   CVG           15.0           0.0
1      1           20           7    MSP   CVG           15.0           0.0
2      1           21           1    MSP   CVG           15.0           0.0
3      1           22           2    MSP   CVG           15.0           0.0
4      1           23           3    MSP   CVG           15.0           0.0
```

```
[13]: # Create dummies for 'ORIGIN' and 'DEST' columns
     # These columns need to be converted into discrete columns containing indicator variables, sometimes known as "dummy" variables.
     # With each column containing 1s and 0s indicating whether a flight originated at the airport that the column represents.
     flight_data = pd.get_dummies(flight_data, columns=['ORIGIN', 'DEST'])
     flight_data.head()
```

```
[13]:   MONTH  DAY_OF_MONTH  DAY_OF_WEEK  CRS_DEP_TIME  ARR_DEL15  ORIGIN_ABE  \
0      1           19           6           15.0           0.0           0
1      1           20           7           15.0           0.0           0
2      1           21           1           15.0           0.0           0
3      1           22           2           15.0           0.0           0
4      1           23           3           15.0           0.0           0
```

	ORIGIN_ABI	ORIGIN_ABQ	ORIGIN_ABR	ORIGIN_ABY	...	DEST_UIN	\
0	0	0	0	0	...	0	
1	0	0	0	0	...	0	
2	0	0	0	0	...	0	
3	0	0	0	0	...	0	
4	0	0	0	0	...	0	

	DEST_USA	DEST_VEL	DEST_VLD	DEST_VPS	DEST_WRG	DEST_XNA	DEST_YAK	\
0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	
2	0	0	0	0	0	0	0	
3	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	

	DEST_YKM	DEST_YUM
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 725 columns]

2 Build Machine Learning Model

```
[14]: import math
flight_data_first_quarter = flight_data[: math.floor(len(flight_data.index)/4)]
flight_data_first_quarter.shape
```

[14]: (159662, 725)

```
[15]: from sklearn.model_selection import train_test_split
# Split DataFrame: flight data into a training set containing 80% of the
# →original data, and a test set containing the remaining 20%
# The random_state parameter seeds the random-number generator used to do the
# →splitting, while the first and second parameters are DataFrames containing the
# →feature columns and the label column.
train_x, test_x, train_y, test_y = train_test_split(flight_data_first_quarter.
# →drop('ARR_DEL15', axis=1), flight_data_first_quarter['ARR_DEL15'], test_size=0.
# →2, random_state=42)
```

```
[16]: train_x.shape
```

[16]: (127729, 724)

```
[17]: train_y.shape
```

[17]: (127729,)

```
[18]: test_x.shape
```

```
[18]: (31933, 724)
```

```
[19]: test_y.shape
```

```
[19]: (31933,)
```

3 Train a classification model

```
[20]: # In this project predicting the probability of a flight will delay, model will
      → be a binary classification model that predicts
      # whether a flight will arrive on-time or late ("binary" because there are only
      → two possible outputs).
      # Use RandomForestClassifier which fits multiple decision trees to the data and
      → uses averaging to boost the overall accuracy and limit overfitting.
      from sklearn.ensemble import RandomForestClassifier
```

```
[21]: from sklearn.metrics import mean_absolute_error
      def get_mae(max_leaf_nodes, n_estimators, train_x, test_x, train_y, test_y):
          model = RandomForestClassifier(max_leaf_nodes=max_leaf_nodes,
          → n_estimators=n_estimators, random_state=3)
          model.fit(train_x, train_y)
          preds_val = model.predict(test_x)
          mae = mean_absolute_error(test_y, preds_val)
          return(mae)
```

```
[31]: # compare MAE with differing values of max_leaf_nodes
      scores = {leaf_size: get_mae(leaf_size, 10, train_x, test_x, train_y, test_y)}
      → for leaf_size in [500, 1000, 2000, 2500, 5000, 10000, 20000, 30000, 50000]}
```

```
[32]: best_tree_size = min(scores, key=scores.get)
      print(best_tree_size)
```

10000

```
[38]: tree_scores = {num_of_trees: get_mae(best_tree_size, num_of_trees, train_x,
      → test_x, train_y, test_y) for num_of_trees in [10, 20, 30, 50, 60, 100]}
```

```
[39]: best_tree_num = min(tree_scores, key=tree_scores.get)
      print(best_tree_num)
```

50

```
[40]: # 'n_estimators' is the number of trees in the forest
      model = RandomForestClassifier(random_state=3, max_leaf_nodes=best_tree_size,
      → n_estimators=best_tree_num)
      model.fit(train_x, train_y)
```

```
[40]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                             max_depth=None, max_features='auto', max_leaf_nodes=10000,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=50, n_jobs=None,
                             oob_score=False, random_state=3, verbose=0, warm_start=False)
```

```
[41]: predicted = model.predict(test_x)
      model.score(test_x, test_y)
```

```
[41]: 0.8210628503429055
```

```
[42]: # One of the best overall measures for a binary classification model is Area
      → Under Receiver Operating Characteristic Curve
      # (sometimes referred to as "ROC AUC"), which essentially quantifies how often
      → the model will make a correct prediction
      # regardless of the outcome.
      # Compute an ROC AUC score for the model

      from sklearn.metrics import roc_auc_score
      probabilities = model.predict_proba(test_x)
```

```
[43]: roc_auc_score(test_y, probabilities[:, 1])
```

```
[43]: 0.7638784856474098
```

```
[44]: # The output from the score method reflects how many of the items in the test
      → set the model predicted correctly.
      # This score is skewed by the fact that the dataset the model was trained and
      → tested with contains
      # many more rows representing on-time arrivals than rows representing late
      → arrivals.
      # Because of this imbalance in the data, it's more likely to be correct if you
      → predict that a flight will be on time than if you predict that a flight will
      → be late.
      # ROC AUC takes this into account and provides a more accurate indication of how
      → likely it is that a prediction of on-time or late will be correct.
```

4 Confusion Matrix

```
[45]: # The confusion matrix quantifies the number of times each answer
      # was classified correctly or incorrectly. Specifically, it quantifies the
      # number of false positives, false negatives, true positives, and true negatives.
```

```
[46]: from sklearn.metrics import confusion_matrix
      confusion_matrix(test_y, predicted)
```

```
[46]: array([[24424,   476],
          [ 5238, 1795]])
```

5 Precision and Recall

```
[47]: from sklearn.metrics import precision_score

train_predictions = model.predict(train_x)
precision_score(train_y, train_predictions)
```

[47]: 0.972072952287901

```
[48]: from sklearn.metrics import recall_score

recall_score(train_y, train_predictions)
```

[48]: 0.42060947683635724

6 Visualize Output of Model


```
[49]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns

sns.set()
```

```
[50]: from sklearn.metrics import roc_curve

fpr, tpr, _ = roc_curve(test_y, probabilities[:, 1])
plt.plot(fpr, tpr)
plt.plot([0, 1], [0, 1], color='grey', lw=1, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

[50]: Text(0, 0.5, 'True Positive Rate')



```
[51]: def insertCode(origin, destination):
    origin = 'ORIGIN_' + origin
    destination = 'DEST_' + destination
    # Create a dict of airports with initial value 0
    airportDict = {}
    for code in train_x.columns.tolist()[4:]:
        airportDict[code] = 0
    airportDict[origin] = 1
    airportDict[destination] = 1
    return airportDict
```

```
[52]: def predict_delay(departure_date_time, origin, destination):
    from datetime import datetime

    try:
        departure_date_time_parsed = datetime.strptime(departure_date_time, '%d/
→%m %H:%M')
    except ValueError as e:
        return 'Error parsing date/time - {}'.format(e)

    month = departure_date_time_parsed.month
    day = departure_date_time_parsed.day
    day_of_week = departure_date_time_parsed.isoweekday()
    hour = departure_date_time_parsed.hour
```



```

origin = origin.upper()
destination = destination.upper()

input = {'MONTH': month,
         'DAY_OF_MONTH': day,
         'DAY_OF_WEEK': day_of_week,
         'CRS_DEP_TIME': hour}
input = [{**input, **insertCode(origin, destination)}]

return model.predict_proba(pd.DataFrame(input))[0][0]

```

```
[53]: predict_delay('1/10 15:45', 'EWR', 'SEA')
```

```
[53]: 0.8414408368109711
```

```
[54]: predict_delay('24/12 09:00', 'LAX', 'EWR')
```

```
[54]: 0.8827298060829265
```


```

[:]: # flights leaving SEA for ATL at 9:00 a.m., noon, 3:00 p.m., 6:00 p.m., and 9:00
      ↪p.m. on January 30 will arrive on time
labels = ('9:00 a.m.', 'Noon', '3:00 p.m.', '6:00 p.m.', '9:00 p.m.')
values = (predict_delay('23/12 09:00', 'EWR', 'LAX'),
          predict_delay('23/12 12:00', 'EWR', 'LAX'),
          predict_delay('23/12 15:45', 'EWR', 'LAX'),
          predict_delay('23/12 18:00', 'EWR', 'LAX'),
          predict_delay('23/12 21:00', 'EWR', 'LAX'))
alabels = np.arange(len(labels))

plt.bar(alabels, values, align='center', alpha=0.5)
plt.xticks(alabels, labels)
plt.ylabel('Probability of On-Time Arrival')
plt.ylim((0.0, 1.0))

```


```
[:]: (0.0, 1.0)
```



```
[ ]: labels = ('Apr 10', 'Apr 11', 'Apr 12', 'Apr 13', 'Apr 14', 'Apr 15', 'Apr 16')
values = (predict_delay('10/04 13:00', 'JFK', 'MSP'),
          predict_delay('11/04 13:00', 'JFK', 'MSP'),
          predict_delay('12/04 13:00', 'JFK', 'MSP'),
          predict_delay('13/04 13:00', 'JFK', 'MSP'),
          predict_delay('14/04 13:00', 'JFK', 'MSP'),
          predict_delay('15/04 13:00', 'JFK', 'MSP'),
          predict_delay('16/04 13:00', 'JFK', 'MSP'))
alabels = np.arange(len(labels))

plt.bar(alabels, values, align='center', alpha=0.5)
plt.xticks(alabels, labels)
plt.ylabel('Probability of On-Time Arrival')
plt.ylim((0.0, 1.0))
```

```
[ ]: (0.0, 1.0)
```



```
[ ]: labels = ('Oct 1', 'Oct 2', 'Oct 3', 'Oct 4', 'Oct 5', 'Oct 6', 'Oct 7')
values = (predict_delay('1/10 21:45', 'JFK', 'ATL'),
          predict_delay('2/10 21:45', 'JFK', 'ATL'),
          predict_delay('3/10 21:45', 'JFK', 'ATL'),
          predict_delay('4/10 21:45', 'JFK', 'ATL'),
          predict_delay('5/10 21:45', 'JFK', 'ATL'),
          predict_delay('6/10 21:45', 'JFK', 'ATL'),
          predict_delay('7/10 21:45', 'JFK', 'ATL'))
alabels = np.arange(len(labels))

plt.bar(alabels, values, align='center', alpha=0.5)
plt.xticks(alabels, labels)
plt.ylabel('Probability of On-Time Arrival')
plt.ylim((0.0, 1.0))
```

```
[ ]: (0.0, 1.0)
```

